Generating NFL Game Headlines from Box Score Statistics

Parker Greene, **Itamar Belson**, and **John Clarke**
{pwgreene, ibelson, jaclarke}@mit.edu

Abstract *In this paper, we introduce a new method by which to extrapolate key game statistics from a complete, purely numerical box score in order to produce a one-sentence summary of the game. Through identifying the particular game figures of greatest importance, the system is able to generate informative headlines that effectively highlight the most significant elements of the particular game. In particular, we describe a system that: (a) successfully utilizes a neural network to identify the significant elements of information necessary to summarize the complete set of data; (b) uses the important data components to generate complete, comprehensive headlines. In this paper we explore the application of the methods to football games played in the National Football League (NFL), but the same approach may be applied to other sports or to a number of other applications.*

I. INTRODUCTION

Sports analytics is an emerging field that seeks to incorporate principles of big data analytics into the world of sports. Data in this arena is extensive, and largely publicly accessible, due in part to the economic appeal of the sports industry as well as the popular public interest in sports. Yet, although considerable work has been completed in the field, much of the work centers on game prediction and individual player evaluation as these applications provide tangible economic benefits to the team franchises. While work for these applications is ongoing, there has been significantly less work in the process of analyzing a complete set of statistics to extrapolate the significant components within a particular game.

Whether a key player's performance or the overall struggle in the contest, concluding such information currently requires either watching the complete game or a comprehensive understanding of and ability to read thorough a publicly available box score. A popular alternative to the two is to read game summaries written by sports professionals, whose expertise in the matter enables them to re-cap a complete game in a condensed form that is understandable to any reader.

This paper explores the potential of machine learning and natural language processing approaches and techniques to extract the important game performances from the box score of NFL games in order to automatically generate these summarizing game headlines. In particular, the paper explores various approaches consisting of neural networks and a variety of generative models to develop an effective system that is capable of completing the currently manual task. Some challenges faced by the system include summarizing a game from a largely naive data set that, in actuality, gives but a glimpse into the complete time-dependent game riddled with the intricacies common to sports. To help combat such issues, we simplified the problem statement to only involve the summarization of events in the actual game, without incorporating information that would require citation outside of the box score. As a result, injury reports, playoff contention, and game disputes are not incorporated into the headlines generated by the developed system.

A. Motivation

From sports to Wall Street, extrapolating key figures from a restricted set of data is an ongoing area of research with applications in a variety of fields. In addition, summarizing the identified key figures in a comprehensive headline may prove instrumental to fields in which language proves more powerful than numbers. As such, although this paper focuses on the information extraction and summarization of football games from box scores, the methods described in this paper may by applied to far reaching applications.

II. Related Work

Generation of somewhat short sentences using Markov chains and stochastic methods have been shown to be quite simple and effective in generating poetry [1] [2], so we focused on applying these same principles in synthesizing valid headlines. This method yields a good way for text generation learned from examples; however, there is relatively much less research in the topic of generation solely from a combination of statistics and their corresponding text examples. Other, more syntactically-based models geared toward dialogue systems [3] [4], we found to be inapplicable to our problem due to the lack of grammatical structure of the headlines in our data. Wen et. al. [5] presented an empirically-tested LSTM model capable of generating linguistically varied responses, but applying a deep, recurrent model to our problem proved hard due to the lack of data and difficulty of integration of the statistics in training.

Football game summaries have been explored by Nichols et. al. [6] and Chakrabarti & Punera [7], but both methods relied on learning data directly from Twitter and were uninformed by the statistics of games, as our model was.

III. Data Collection

The data used by the system was extracted directly from ESPN's publicly accessible website1. The website contains comprehensive historical statistics for the past ten seasons of NFL games (2007-2016), with a complete box score for each game played throughout each season. From these box score, we extracted twenty-eight key statistics that would provide substantial insight into each overall game. These key statistics are summarized in Figure —-. In addition to the statistics extracted directly from the box score, we also incorporated team specific information for each game, such as the teams' city and acronym. This information proved useful in the annotation of the game headlines described below.

In addition to the box score data, each game on the website also has an accompanying article headline that was manually written by a professional sports writer following the game. We scraped this headline for each of the games for use as training data for our system. However, as described in the next section, our approach took into consideration not the headline provided directly from the website, but a general form of the headline. To generalize the headlines, we removed all game specific names and figures and replaced them with identifying tags. A few examples of the general form headlines can be seen in Figure —-. The initial process was completed automatically but required additional manual processing to ensure accurate annotation of the data. After the data was compiled and annotated, we removed any game data with an accompanying headline that incorporated information external to the actual gameplay, such as injury reports, playoff contention, and game disputes. In conclusion, our final dataset consisted of data and a headline from 1742 individual games.

IV. Approach

A. Statistic Feature Vectors

B. Neural Net

The NN accepts a vector x of statistics and returns a vector containing the probabilities of words existing in the headline associated with those statistics. We explain our choice of statistic representations in section V. First, projects the statistics vector x through two hidden layers of size $2|x|$. The hidden layer outputs are activated by a Rectified Linear function (ReLU). Then, the result is project to a vector space of size $|V|$, where V is the vocabulary including start and end characters, finally activated by a softmax function. Refer to 2 for an illustration of the NN architecture. We initialize the weights by sampling from a Gaussian distribution $N(0, 1)$. To train the model, we optimize a categorical cross-entropy loss function using stochastic gradient descent (SGD) with adaptive maximum descent (Adam). This training method suits our objective of producing multi-class predictions. Lastly, we regularize our model on two fronts: (1) with a learning rate parameter and (2) neural dropout, where we ignore the weights connected to neurons that are "turned off" with probability p during training. With our vector y of word probabilities $p(w_i)$ for $w_i \in V$, we move on to trigram MC.

C. Re-weighted Trigram Markov Generator

We initialize the MC with transition probabilities estimated from the corpus of 1742 game headlines. For this task, we simply calculated the
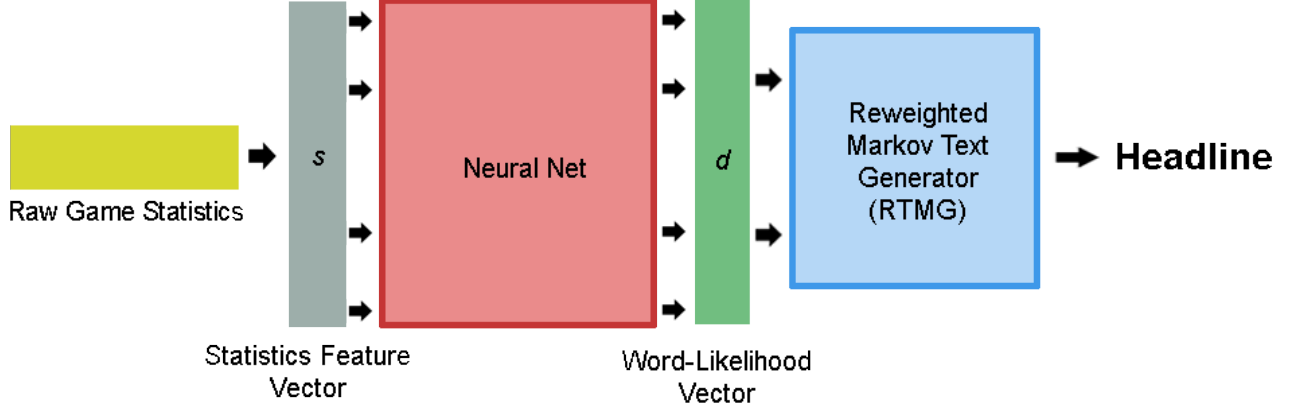
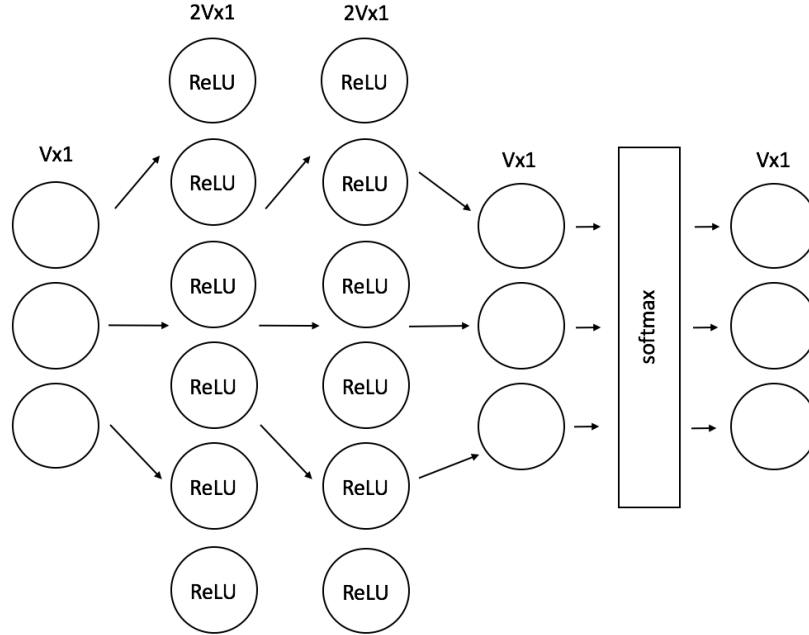Fig. 1. Box Diagram of the full headline generation system

Fig. 2. Overview of the neural network model. Input to the network is the feature vector s and output is the word-likelihood vector d.

maximum likelihood estimates for each word as $p(w_i|w_{i-1}, w_{i-2}) = \frac{count(w_i, w_{i-1}, w_{i-2})}{count(w_{i-1}, w_{i-2})}$, yielding a matrix of transition probabilities, T. Refer to Figure 2 for an example illustration of the MC model. At this point, the MC could be used to create reasonably semantically and syntactically correct sentences. The model takes an element-wise product of our word probability vector d from (1) and each row of T, which describe the transitions $u, v \rightarrow w$, for all words $w \in V$ given words u, v. These probabilities serve as discounting factors to properly

weigh down the transition probabilities to words unlikely to appear in the headline. Leveraging this information results in generally grammatical sentences that are generated around the relevant vocabulary. Code for the re-weighting algorithm in listed in 1. The function takes as input T, the existing transition matrix as estimated from the corpus of headlines. T is re-weighted by d, which is the output of the neural net and has the same length as each row of T. The re-weighted transition matrix, M is returned. Note that d changes based on s, the input

feature vector to the neural net, so T is re-weighted only to generate sentences for a particular s and then re-weighted differently for a different s.

Sentences are generated by doing a random walk over the possible states, where each word, including the start and end symbols, are considered states. The generator starts at a pair of two start symbols and then randomly walks over the words according to the transition probabilities until the end symbol is reached. If the current state is (u, v), then the next state is chosen according to the probabilities in M corresponding to pair (u, v).

```python
def reweight_markov(d, T):
  M = copy(T)
  for u in range(0, n*n):
    for v in range(0, n):
      M[u][v] = T[u][v] * d[v]
  return M
```

Listing 1: Re-weighting the Markov text generator

MARKOV MODEL ARCHITECTURE **??**.

V. RESULTS

VI. CONTRIBUTION

A. Parker Greene

I started by working on the model design, and after a lot of research into text generation, I initially came up with the idea of not using the raw statistics as input directly to the sentence generation model, but instead separating the system into the two separate models: one for creating an embedding from the statistics and one for using this embedding to generate the headlines. After the initial model design, I focused on the problem of text generation. The initial approach I considered was the bigram Markov chain, which had the ability to generate sentences, but could only learn directly from the corpus of headlines and thus relied on the neural net output. I also explored the possibility of using both word-level and text-level sequence generation using an LSTM network. This approach worked well for learning from the corpus how to generate sentences, but unlike the Markov model, it could not easily be integrated with the existing embeddings output by the neural net. I also explored the possibility of using PCFG's to generate the headlines, which would have been able to assure the headlines had a given structure, but, as was touched upon in section III, I found they could not properly work with the headlines we had, since the headlines did not follow proper grammar. We contributed equal work to this write-up and the presentation.

B. Itamar Belson

C. John Clarke

REFERENCES

[1] Kenner, Hugh; O'Rourke, Joseph (November 1984). "A Travesty Generator for Micros". BYTE. 9 (12): 129-131, 449-469.

[2] Hartman, Charles (1996). *Virtual Muse: Experiments in Computer Poetry*. Hanover, NH: Wesleyan University Press.

[3] Alice H. Oh and Alexander I. Rudnicky. 2000. "Stochastic language generation for spoken dialogue systems". In Proceedings of the 2000 ANLP/NAACL Workshop on Conversational Systems - Volume 3, ANLP/NAACL-ConvSyst '00.

[4] Adwait Ratnaparkhi. 2002. "Trainable approaches to surface natural language generation and their application to conversational dialog systems". Computer Speech and Language.

[5] Tsung-Hsien Wen, Milica Gasic, Dongho Kim, Nikola Mrksic, Pei-Hao Su, David Vandyke, and Steve Young. 2015. "Stochastic language generation in dialogue using recurrent neural networks with convolutional sentence reranking". In Proceedings of SIG-dial. Association for Computational Linguistics.

[6] Jeffrey Nichols, Jalal Mahmud, Clemens Drews. 2012. "Summarizing Sporting Events Using Twitter". IBM Research - Almaden

[7] Chakrabarti, D., and Punera, K. 2011. "Event summarization using tweets". Proceedings of the Fifth International AAAI Conference on Weblogs and Social Media.